

# LangChain 3-Hour Exercise Guide: News API + RAG System

## OVERVIEW

**Duration:** 3 Hours Total (Exercise 1: 1.5 hours | Exercise 2: 1.5 hours)

**Goal:** Master LangChain loaders and complete RAG pipeline using free APIs

---

## EXERCISE 1: News API Loader with AI Enhancement (90 minutes)

### FREE NEWS API SETUP

#### Using NewsData.io API (Recommended - Like OpenWeather)

- **Sign up:** Go to <https://newsdata.io/register> (Free account)
- **Free Tier:** 200 requests/day (perfect for learning)
- **API Pattern:** Exactly like your weather API example

Base URL: `https://newsdata.io/api/1/news`

Required Parameter: `apikey=YOUR_API_KEY`

Optional Parameters: `country=us, language=en, category=general`

#### Example API Call Structure (Same as Weather API):

```
https://newsdata.io/api/1/news?  
apikey=YOUR_API_KEY&country=us&language=en
```

## STEP-BY-STEP IMPLEMENTATION

### Step 1: Environment Setup (10 minutes)

1. Create `.env` file and add your NewsData API key
2. Create `config.py` with API key loader function (copy from your weather example)
3. Test basic API connection using `requests.get()`

### Step 2: Create Custom NewsLoader Class (20 minutes)

**Follow this exact pattern from your weather example:** - Create class similar to `WeatherAPILoader` - Add constructor with `api_key` parameter - Add `load()` method that makes API request - Return the JSON response (list of news articles)

### Step 3: AI Enhancement Pipeline (30 minutes)

**Create AI enhancement for each news article:** - Use your existing PromptTemplate pattern - Create prompts that take news article and generate: - AI summary (2 sentences max) - Sentiment analysis (positive/negative/neutral) - Key topics extraction (3-5 topics) - Credibility assessment (brief explanation)

### Step 4: JSON Output Generator (20 minutes)

**Combine original news + AI enhancements:** - Loop through news articles - For each article, get AI enhancements - Structure final JSON with original data + AI additions - Save to enhanced\_news.json file

### Step 5: Testing & Validation (10 minutes)

- Test with 5-10 articles
- Verify JSON structure
- Check AI enhancement quality

**Expected Output:** enhanced\_news.json file with original news data + AI enhancements

---

## EXERCISE 2: Organization RAG System (90 minutes)

### \*\* FOCUS: Complete RAG Pipeline\*\*

**Core Learning:** Data → Chunks → Vector Store → Retrieval → LLM Response

### \*\* STEP-BY-STEP RAG IMPLEMENTATION\*\*

#### Step 1: Choose Organization & Data Sources (15 minutes)

**Pick ONE organization with rich free data:** - Tesla (recommended - lots of public data) - NASA - Your university - Local government - Microsoft

**Data Sources to Use:** - Official website pages (3-5 key pages) - Wikipedia page about the organization - 1-2 PDF documents (annual report, company overview)

#### Step 2: Data Loading Phase (20 minutes)

**Use LangChain loaders to collect ALL data:**

**Web Content:** - Use WebBaseLoader for official website pages - Load 3-5 important pages (About, Products, History, etc.)

**Wikipedia:** - Use WikipediaLoader to get comprehensive organization info - Load main article + related pages

**PDF Documents:** - Use PyPDFLoader for annual reports or public documents - Download PDFs first, then load with PyPDFLoader

**Combine all documents into one list**

### **Step 3: Chunking Strategy (15 minutes)**

**Split all documents for better retrieval:** - Use RecursiveCharacterTextSplitter - Set chunk\_size=1000, chunk\_overlap=200 - Apply to ALL loaded documents - Count total chunks created

### **Step 4: Vector Store Creation (15 minutes)**

**Store chunks in Chroma database:** - Use your existing embeddings setup - Create Chroma database with persist\_directory - Store ALL document chunks - Name database after your organization (e.g., ./chroma\_tesla)

### **Step 5: Boundary Enforcement System (10 minutes)**

**Create organization-only question filter:** - Create simple function to check if question mentions your organization - Use basic keyword matching or simple LLM prompt - Return "I only answer questions about [ORGANIZATION]" for off-topic queries

### **Step 6: Retrieval System (10 minutes)**

**Set up document retrieval:** - Create retriever from your vector store - Set k=3 (retrieve top 3 relevant chunks) - Test retrieval with sample questions about your organization

### **Step 7: Complete RAG Pipeline (5 minutes)**

**Chain everything together:** 1. Check if question is about organization 2. If yes: retrieve relevant documents 3. Format documents as context 4. Create prompt with context + question 5. Get LLM response 6. Return answer with source indication

## **COMMAND LINE INTERFACE REQUIREMENTS**

**Your CLI must handle:**

Welcome to [ORGANIZATION] Assistant  
Ask questions about [ORGANIZATION] only!

User: "What does Tesla do?"

Assistant: [Answer based on retrieved documents]

User: "What is Apple's stock price?"

Assistant: "I only answer questions about Tesla. Please ask about Tesla."

User: "quit"

System: Exit gracefully

## TESTING YOUR RAG SYSTEM

### Test Cases to Verify:

1. **Boundary Enforcement:** Ask about different organizations
2. **Document Retrieval:** Ask specific questions that should find relevant chunks
3. **Source Context:** Verify answers use retrieved information
4. **Edge Cases:** Empty questions, very long questions

### Quality Checks:

- Can retrieve information from web pages? ✓
  - Can retrieve information from Wikipedia? ✓
  - Can retrieve information from PDFs? ✓
  - Rejects off-topic questions? ✓
  - Provides relevant answers with context? ✓
- 

## DELIVERABLES CHECKLIST

### Exercise 1 Outputs:

- ☐ Working NewsLoader class
- ☐ AI enhancement pipeline
- ☐ enhanced\_news.json file with 10+ articles
- ☐ Clean, documented Python code

### Exercise 2 Outputs:

- ☐ Multi-source data loading (web + wiki + PDF)
  - ☐ Chunked and stored documents in Chroma
  - ☐ Working CLI with organization boundary enforcement
  - ☐ Successful retrieval and response generation
-

# TIME MANAGEMENT GUIDE

## Hour 1: News API Exercise

- 0-10 min: API setup and testing
- 10-30 min: NewsLoader class creation
- 30-60 min: AI enhancement pipeline

## Hour 2: Complete News Exercise + Start RAG

- 0-30 min: JSON output and testing (complete Exercise 1)
- 30-45 min: Choose organization and plan data sources
- 45-60 min: Load all data sources

## Hour 3 for DEEPSEED's WAR: Complete RAG System

- 0-15 min: Chunk and store in vector database
- 15-30 min: Build retrieval system
- 30-45 min: Create boundary enforcement and CLI
- 45-60 min: Testing and debugging

---

## COMMON PITFALLS TO AVOID

1. **API Limits:** Don't make too many API calls during testing
  2. **Large Files:** Check PDF size before loading (keep under 10MB)
  3. **Chunk Size:** If retrieval is poor, adjust chunk size
  4. **Organization Scope:** Keep boundary checking simple but effective
  5. **Error Handling:** Add basic try-catch for API calls
- 

## SUCCESS CRITERIA

**Exercise 1 Complete When:** - News API returns data successfully - AI enhancements work for each article - JSON file contains structured, enhanced news data

**Exercise 2 Complete When:** - Can load data from multiple sources - Documents are chunked and stored in Chroma - CLI answers organization questions correctly - Rejects off-topic questions appropriately - Retrieval system finds relevant information

**Both exercises should run without errors and produce expected outputs.**